



Cardiovascular Disease Risk Prediction Model

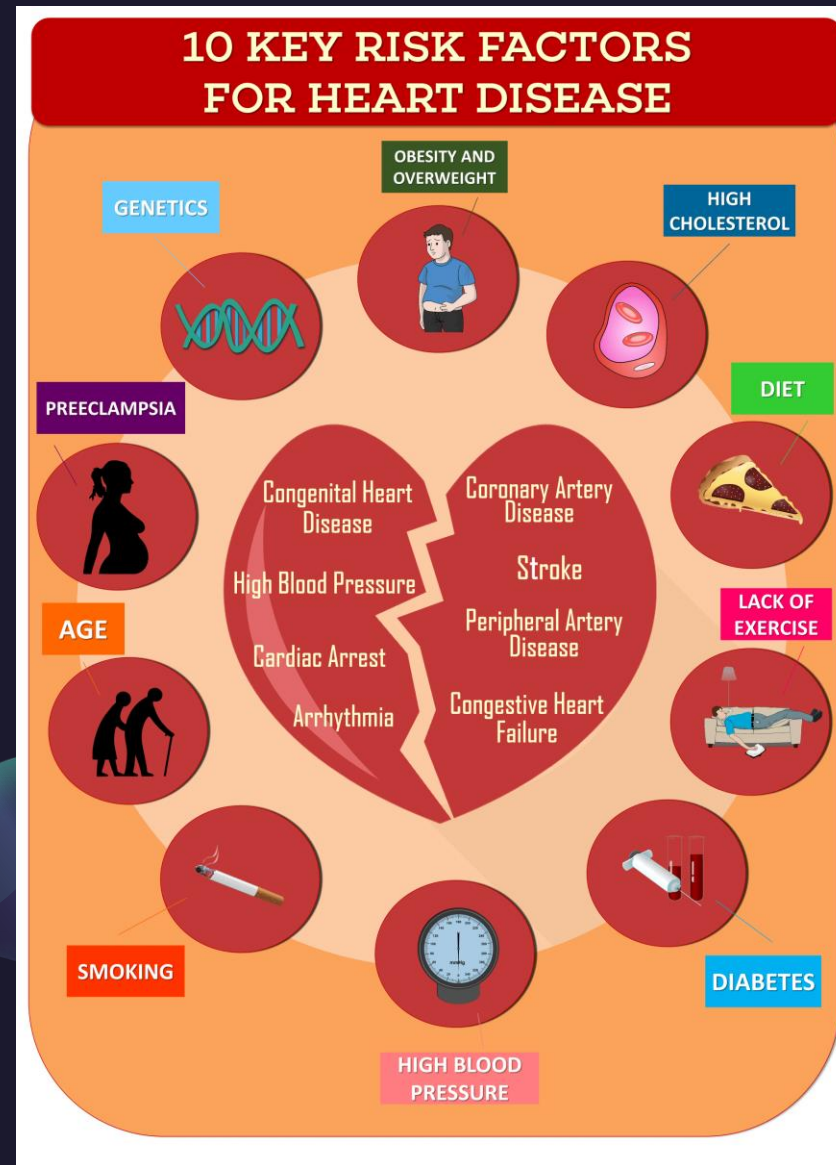
Presentation by: Aline Vo, Annie Joseph Rajan,
Stuti Poudel, Brian Stumm, and Jules Gikundiro

Heart Disease

- The leading cause of death globally ~ 18Million/Year
- Cost the US ~ \$240 Billion
- Goal – Risk Prediction

Risk Factors:

- Smoking
- Exercise
- Cancer
- General Health
- BMI
- Depression
- Diabetes





Data Cleaning and Preprocessing

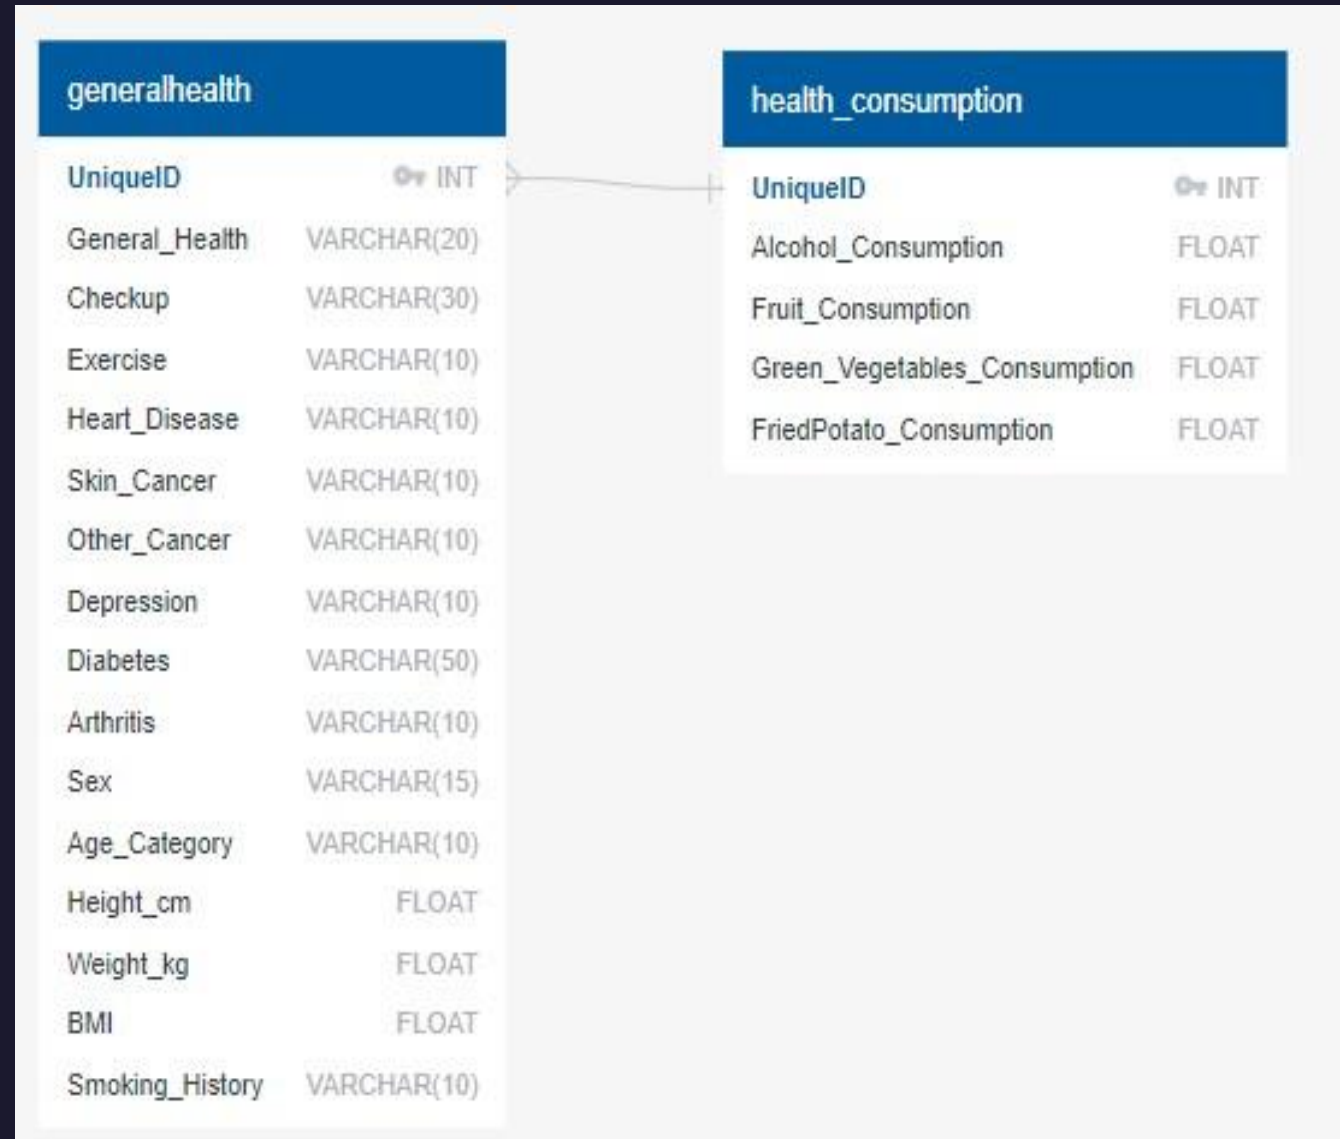
Data Cleaning

- Created a unique ID for each row of information
- Created two DataFrames
 - Consumption DB
 - Non-Consumption DB
- Created CSV files
- Imported clean data into SQL



SQL

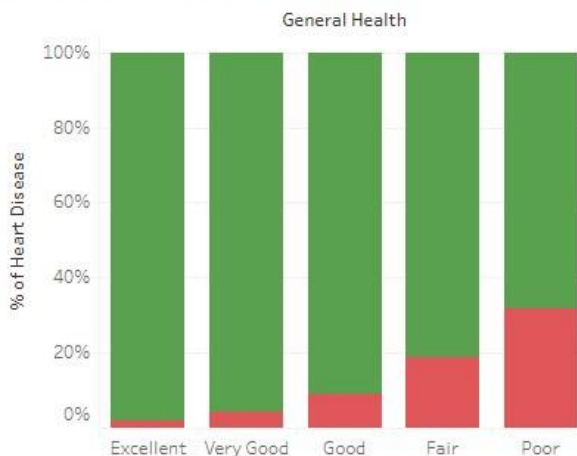
- Created two tables
 - General Health
 - Health consumption
- ERD



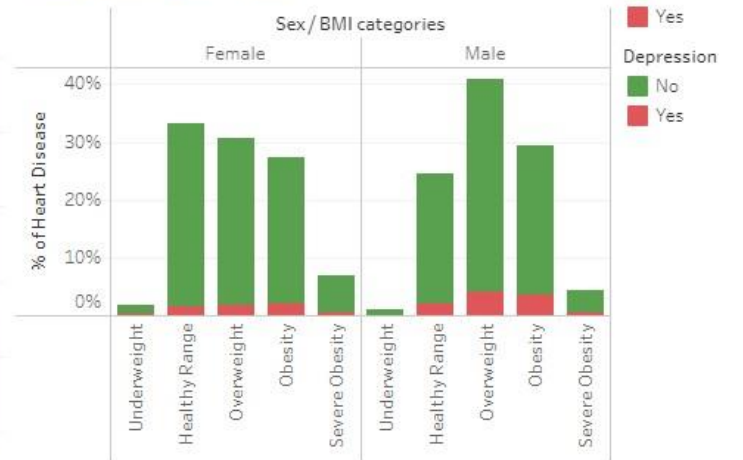
Visualization Dashboard

Factors that contribute to Heart Disease


Heart Disease - General Health




Heart Disease - BMI



Depression relates to Heart Disease



Diabetes related to Heart Disease

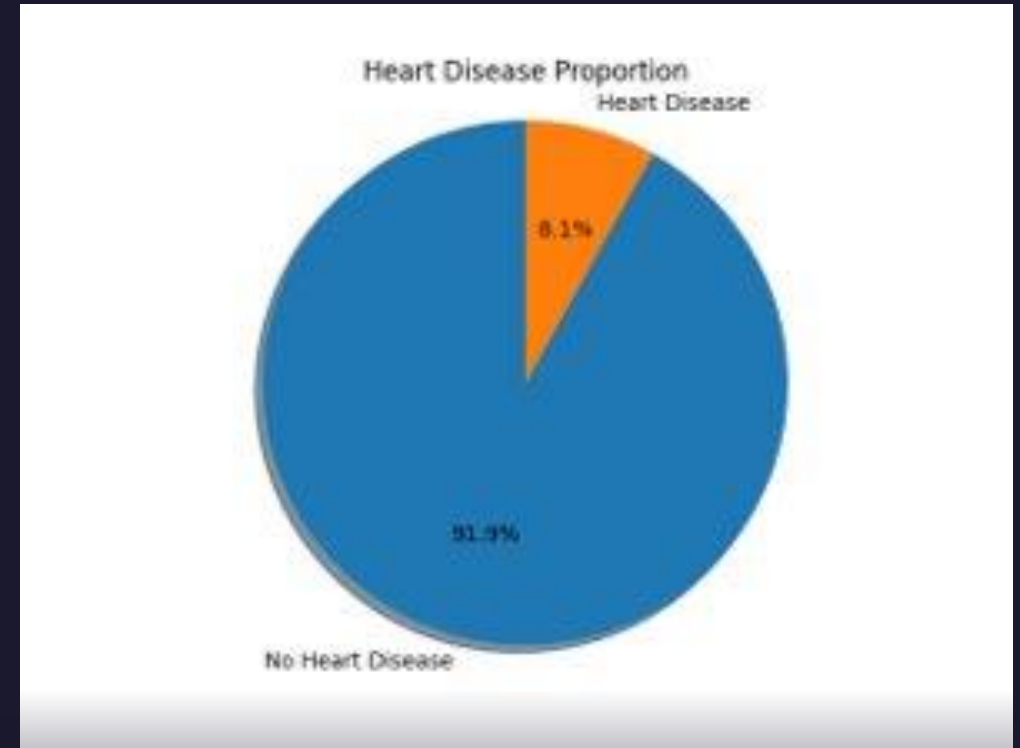



Machine Learning Models

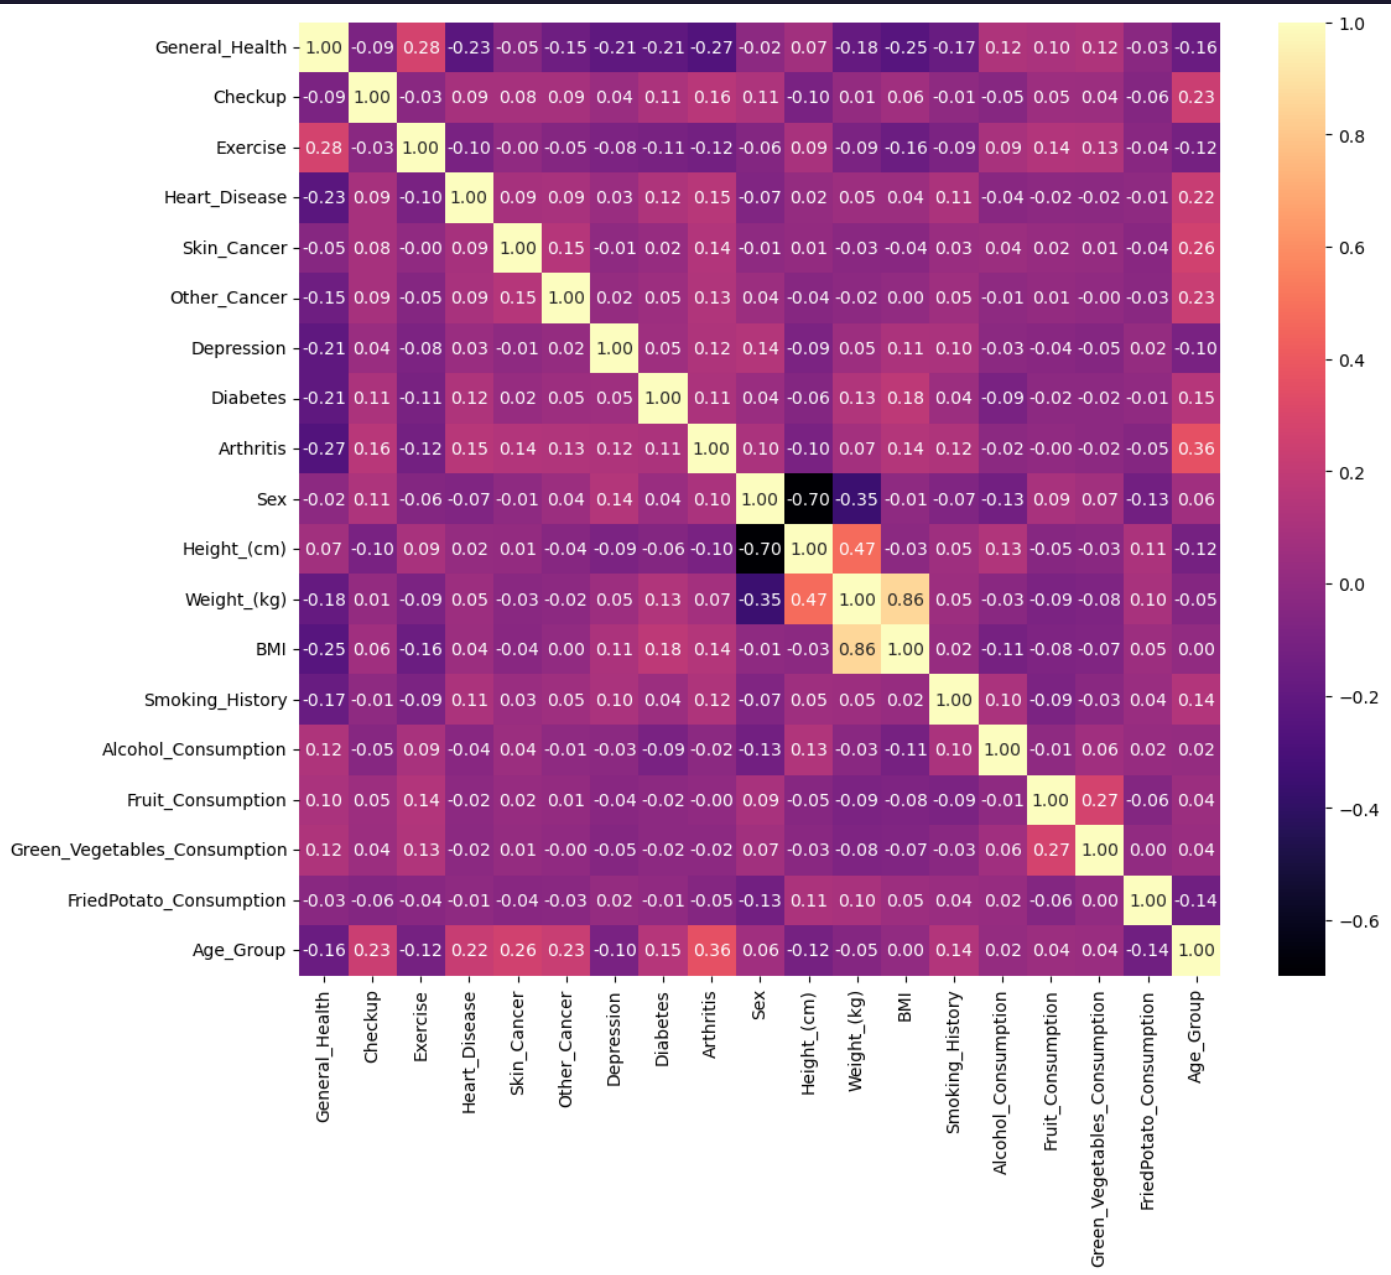
Neural Network Model, Support Vector Machines Model, and Random Forest Classifier Model

Preprocessing

- Examined distribution – imbalanced data
- Oversampling methods



Correlation Analysis



Neural Network Model

Input Layer :

- The number of neurons equal to the number of features in the input data.

First Hidden Layer:

- Number of Neurons: 80
- Activation Function: ReLU (Rectified Linear Unit)

Second Hidden Layer:

- Number of Neurons: 30
- Activation Function: ReLU

Output Layer:

- Number of Neurons: 1 (since it's a binary classification task)
- Activation Function: Sigmoid (to output probabilities for binary classification)

Neural Network Model Final Results

Before Optimization

```
2413/2413 - 2s - 994us/step - accuracy: 0.9180 - loss: 0.2314  
Loss: 0.2313869297504425, Accuracy: 0.9180330038070679
```

After First Optimization

```
2413/2413 - 2s - 971us/step - accuracy: 0.9183 - loss: 0.2329  
Loss: 0.23291058838367462, Accuracy: 0.9182920455932617
```

After Second Optimization

```
2413/2413 - 2s - 848us/step - accuracy: 0.9203 - loss: 0.2224  
Loss: 0.2224111258983612, Accuracy: 0.9202605485916138
```



Support Vector Machines Model (SVM)

SVM Model

```
from sklearn.svm import SVC

# Create the SVM model with a rbf kernel
model = SVC(kernel='rbf')

# Fit the model to your training data
model.fit(X_train, y_train)
```

SVC

SVC ()

```
# Calculate the classification report
from sklearn.metrics import classification_report
predictions = model.predict(X_test)
print(classification_report(y_test, predictions))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.92 | 1.00 | 0.96 | 70955 |
| 1 | 0.00 | 0.00 | 0.00 | 6259 |
| accuracy | | | 0.92 | 77214 |
| macro avg | 0.46 | 0.50 | 0.48 | 77214 |
| weighted avg | 0.84 | 0.92 | 0.88 | 77214 |

SVM Model with SMOTE

```
# Applying SMOTE
sm = SMOTE(random_state=42)
X_train_res, y_train_res = sm.fit_resample(X_train, y_train)

# Now, we can use the resampled data to train your model
model = SVC(kernel='rbf')
model.fit(X_train_res, y_train_res)
```

Accuracy: 0.7757142487113736

Confusion Matrix:
[[56205 14750]
[2568 3691]]

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.96 | 0.79 | 0.87 | 70955 |
| 1 | 0.20 | 0.59 | 0.30 | 6259 |
| accuracy | | | 0.78 | 77214 |
| macro avg | 0.58 | 0.69 | 0.58 | 77214 |
| weighted avg | 0.90 | 0.78 | 0.82 | 77214 |

SVM with Balanced class weight

```
from sklearn.svm import SVC

# Create an SVC model with balanced class weights
model = SVC(kernel='rbf', class_weight='balanced')

# Train the model with your data
model.fit(X_train, y_train)
```

SVC

SVC (class_weight='balanced')

Test Acc: 0.712

```
# Calculate the classification report
from sklearn.metrics import classification_report
predictions = model.predict(X_test)
print(classification_report(y_test, predictions))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.98 | 0.70 | 0.82 | 70915 |
| 1 | 0.20 | 0.81 | 0.32 | 6299 |
| accuracy | | | 0.71 | 77214 |
| macro avg | 0.59 | 0.76 | 0.57 | 77214 |
| weighted avg | 0.91 | 0.71 | 0.78 | 77214 |



SVM Model Final Report



```
from sklearn.svm import SVC

# Manually specifying the class weights
# Giving class 1 a higher weight
class_weights = {0: 1, 1: 10}

# Create an SVC model with custom class weights
model = SVC(kernel='rbf', class_weight=class_weights)
```



Test Acc: 0.739

```
# Calculate the classification report
from sklearn.metrics import classification_report
predictions = model.predict(X_test)
print(classification_report(y_test, predictions))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.97 | 0.74 | 0.84 | 70915 |
| 1 | 0.21 | 0.78 | 0.33 | 6299 |
| accuracy | | | 0.74 | 77214 |
| macro avg | 0.59 | 0.76 | 0.58 | 77214 |
| weighted avg | 0.91 | 0.74 | 0.80 | 77214 |

Random Forest Classifier

```
# Dependencies
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, r2_score
```

```
# Instantiate a Random Forest Classifier model
rf_model = RandomForestClassifier(n_estimators=100, class_weight='balanced', random_state=1)
```

```
# Fit the model with training data
rf_model.fit(X_train, y_train)
```

```
RandomForestClassifier
RandomForestClassifier(class_weight='balanced', random_state=1)
```

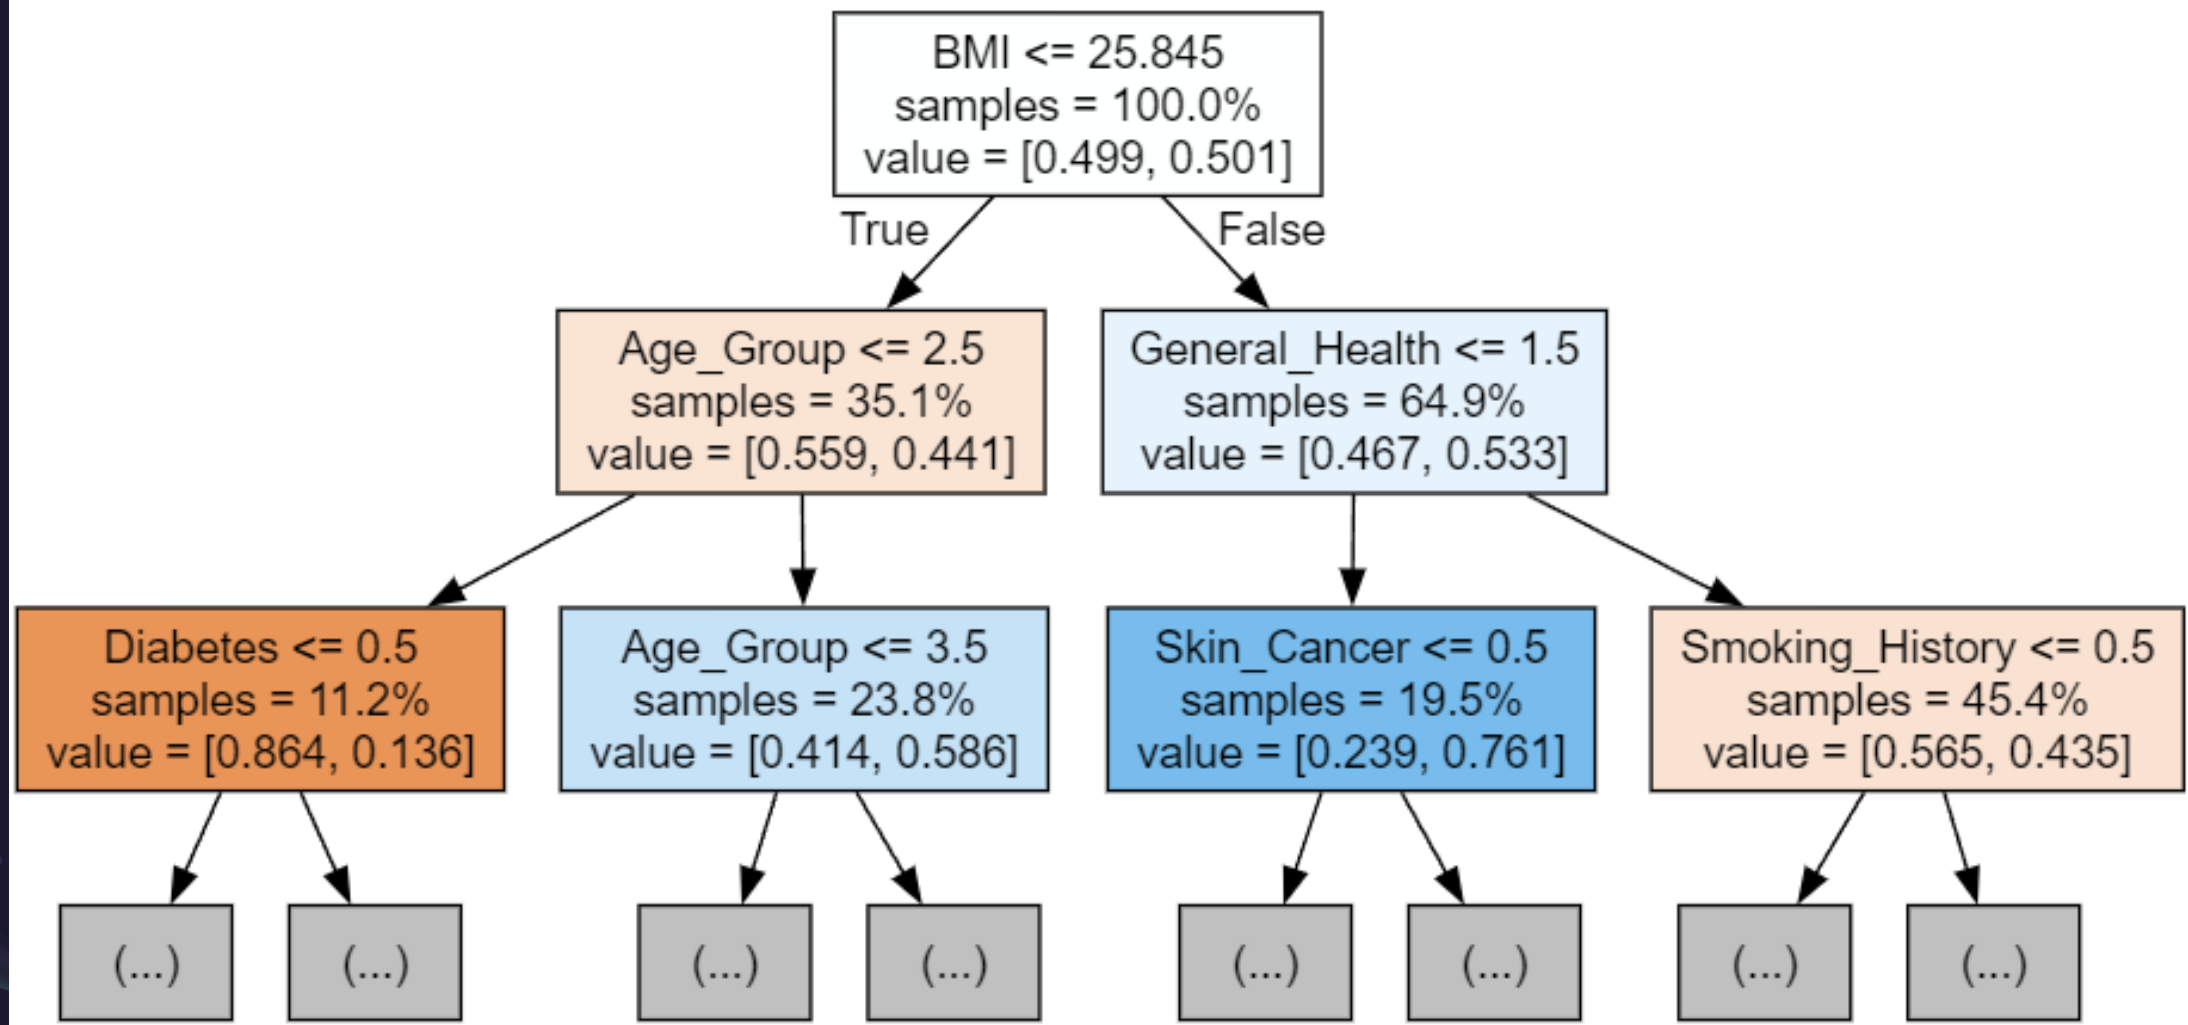
```
# Make predictions on the test set
predictions = rf_model.predict(X_test)
```

Confusion Matrix

```
# Generate a confusion matrix for the model
display(cm_df)
```

| | Predicted 0 | Predicted 1 |
|----------|-------------|-------------|
| Actual 0 | 54823 | 1999 |
| Actual 1 | 7 | 56725 |

Decision Tree from the Forest



Random Forest Classifier Final Report



```
# Look at the accuracy score
print(f"Accuracy Score : {acc_score}")
```

Accuracy Score : 0.9825369427761241

```
# Look at the classification report
print("Classification Report")
print(classification_report(y_test, predictions))
```

| Classification Report | | | | |
|-----------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 1.00 | 0.97 | 0.98 | 56822 |
| 1 | 0.97 | 1.00 | 0.98 | 56732 |
| accuracy | | | 0.98 | 113554 |
| macro avg | 0.98 | 0.98 | 0.98 | 113554 |
| weighted avg | 0.98 | 0.98 | 0.98 | 113554 |

```
... Classification Report
      precision    recall  f1-score   support


0         0.93      0.91      0.92     56822
1         0.91      0.93      0.92     56732

accuracy          0.92     113554
macro avg         0.92      0.92      0.92     113554
weighted avg      0.92      0.92      0.92     113554
```

Opt 2: SMOTE

| Classification Report | | | | |
|-----------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 0.94 | 0.92 | 0.93 | 56822 |
| 1 | 0.92 | 0.94 | 0.93 | 56732 |
| accuracy | | | 0.93 | 113554 |
| macro avg | 0.93 | 0.93 | 0.93 | 113554 |
| weighted avg | 0.93 | 0.93 | 0.93 | 113554 |

Opt 3: BorderlineSMOTE

 Model: RandomOverSampler

Challenges and Limitations

- Feature Correlation Analysis shows weak to no correlation to target variable
- Missing Key Features such as, High Blood Pressure, High Cholesterol, Stress, and Family history
- Complexity of the Dataset biased toward minority classes
- Inadequate domain knowledge
- More binary variables than numerical variables



Conclusion



- Random Forest Classifier Model with RandomOverSampler had the best performance
- Accuracy score: 98.23%
- Confusion Matrix
 - True Positive (Actual 1) 56,725 times
 - True Negative (Actual 0) 54,823 times
 - False Negative (Actual 1) 7
 - False Positive (Actual 0) 1999



```
# Look at the accuracy score  
print(f"Accuracy Score : {acc_score}")
```

Accuracy Score : 0.9823343959701992

```
# Look at the classification report  
print("Classification Report")  
print(classification_report(y_test, predictions))
```

| Classification Report | | | | |
|-----------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 1.00 | 0.96 | 0.98 | 56822 |
| 1 | 0.97 | 1.00 | 0.98 | 56732 |
| accuracy | | | 0.98 | 113554 |
| macro avg | 0.98 | 0.98 | 0.98 | 113554 |
| weighted avg | 0.98 | 0.98 | 0.98 | 113554 |

```
# Generate a confusion matrix for the model  
display(cm_df)
```

| | Predicted 0 | Predicted 1 |
|----------|-------------|-------------|
| Actual 0 | 54823 | 1999 |
| Actual 1 | 7 | 56725 |

Thank you

Q & A

